

```

1  # -*- coding: utf-8 -*-
2  """Logistic_Regression
3
4  Automatically generated by Colaboratory.
5
6  Original file is located at
7      https://colab.research.google.com/drive/1qioJbplkpgPKdiDEP2SYKmu6t7V-Maoo
8
9  <center><h1>Mini Project 1 - Logistic Regression</h1>
10 <h4>The hyperparameters and models used in this file are chosen based on the
    findings in the testing file.</h4></center>
11
12 <h3>Team Members:</h3>
13 <center>
14 Yi Zhu, 260716006<br>
15 Fei Peng, 260712440<br>
16 Yukai Zhang, 260710915
17 </center>
18 """
19
20 from google.colab import drive
21 drive.mount('/content/drive')
22
23 import numpy as np
24 import pandas as pd
25 import matplotlib.pyplot as plt
26 import seaborn as sns
27
28 path1 = "/content/drive/My Drive/ECSE_551_F_2020/Mini_Project_01/hepatitis.csv"
29 path2 = "/content/drive/My Drive/ECSE_551_F_2020/Mini_Project_01/bankrupcy.csv"
30
31 class LogisticRegression:
32     '''
33         This is the logistic regression class, containing fit, perdict and
34         accu_eval functions,
35         as well as many other useful functions.
36     '''
37     def __init__(self, data, folds, lr=0.01, max_iter=10000, beta=0.99, epsilon
38 =5e-3):
39         self.data = data
40         self.folds = folds
41         self.lr = lr
42         self.max_iter = max_iter
43         self.beta = beta
44         self.epsilon = epsilon
45     def shuffle_data(self):
46         '''
47             This function randomly shuffles the input dataset.
48         '''
49         # Load data from data file.
50         self.data.insert(0, column='Bias', value=1)
51         self.data = self.data.sample(frac=1)
52     def partition(self, fold):
53         '''
54             This function divides the dataset into training and validation set.
55         '''
56         fold - the current fold
57         data = self.data
58         # to exclude last term in previous partition for training data

```

```

61     train_add = 1 if fold < self.folds else 0
62     # to exclude last term in previous partition for testing data
63     test_add = 1 if fold > 0 else 0
64
65     # number of data sets
66     n = len(self.data)
67
68     train_set_1 = data.iloc[0:int((fold)/self.folds*n), :]
69     train_set_2 = data.iloc[int((fold+1)/self.folds*n)+train_add:n, :]
70     train_set = pd.concat([train_set_1, train_set_2])
71
72     test_set = data.iloc[int((fold)/self.folds*n+test_add):int((fold+1)/
self.folds*n), :]
73
74     train_X = train_set.iloc[:, :-1].values
75     train_y = train_set.iloc[:, -1].values
76     train_y = np.reshape(train_y, (-1,1))
77
78     test_X = test_set.iloc[:, :-1].values
79     test_y = test_set.iloc[:, -1].values
80     test_y = np.reshape(test_y, (-1,1))
81
82     return train_X, train_y, test_X, test_y
83
84     def normalization(self, X, v_X):
85         '''
86             This function performs the z-score normalization
87
88             X - training data
89             v_X - validation data
90         '''
91         mean = np.mean(X[:,1:], axis = 0)
92         sigma = np.std(X[:,1:], axis = 0)
93         mean = np.reshape(mean, (1,-1))
94         sigma = np.reshape(sigma, (1,-1))
95         X[:,1:] = (X[:,1:] - mean) / sigma
96         v_X[:,1:] = (v_X[:,1:] - mean) / sigma
97         return X, v_X
98
99     def fit(self, X, y, v_X, v_y, normalize=False):
100         '''
101             This function takes the training data X and its corresponding
labels vector y
102             as well as other hyperparameters (such as learning rate) as input,
103             and execute the model training through modifying the model
parameters (i.e. W).
104
105             X - training data
106             y - class of training data
107             v_X - validation data
108             v_y - class of validation data
109             epsilon - the threshold value for gradient descent
110             normalize - whether to perform normalization
111         '''
112         gradient_values, t_acc_val, v_acc_val = [], [], []
113
114         if normalize:
115             X, v_X = self.normalization(X, v_X)
116
117         # Retrive the learning rate, maximum iteration, momentum (beta)
118         lr, max_iter, beta, epsilon = self.lr, self.max_iter, self.beta, self.
epsilon
119

```

```

120     # initial weight vector
121     w = np.zeros((len(X[0]), 1))
122     # record the best weight vector
123     best_w = w
124     # iteration number, validation accuracy, last validation accuracy
125     # the step to take in gradient descent, maximum validation accuracy
126     iteration, v_acc, step, v_acc_max = 0, 0, 0, 0
127
128     dw = np.inf
129     # if the gradient delta w is smaller than threshold or achieved
    maximum iteration, stop
130     while (np.linalg.norm(dw) > epsilon and iteration <= max_iter):
131         dw = self.gradient(X, y, w)
132         gradient_values.append(np.linalg.norm(dw))
133         # if beta = 0, it will be the same as general gradient descent
134         step = beta * step + (1 - beta) * dw # gradient descent with
    momentum
135         w = w - lr * step
136
137         # predict once every 10 iterations
138         if iteration % 10 == 0:
139             t_y_pred = self.predict(X, w)
140             t_acc = self.accu_eval(t_y_pred, y)
141             v_y_pred = self.predict(v_X, w)
142             v_acc = self.accu_eval(v_y_pred, v_y)
143
144             # record the next best value
145             if v_acc >= v_acc_max:
146                 v_acc_max = v_acc
147                 best_w = w
148                 self.marker = iteration # move the iteration marker
149
150             t_acc_val.append(t_acc)
151             v_acc_val.append(v_acc)
152
153             iteration = iteration + 1
154
155     return gradient_values, t_acc_val, v_acc_val, best_w
156
157     def predict(self, X, w):
158         '''
159         This function takes a set of data as input and outputs predicted
    labels for the input points.
160         '''
161         result = self.log_func(np.dot(X, w))
162         # the prediction result converted to binary
163         predict_bin = []
164         for i in result:
165             if i >= 0.5:
166                 predict_bin.append(1)
167             else:
168                 predict_bin.append(0)
169         return predict_bin
170
171     def accu_eval(self, y_pred, y):
172         '''
173         This function evaluates the models' accuracy.
174         '''
175         count = 0
176         for i in range(len(y_pred)):
177             if y_pred[i] == y[i]:
178                 count = count + 1
179         # return the accuracy ratio: #corret prediction / #data points

```

```

180         return count / len(y)
181
182     def log_func(self, alpha):
183         return 1 / (1 + np.exp(-alpha))
184
185     def gradient(self, X, y, w):
186         N = len(X[0])
187         y_hat = self.log_func(np.dot(X, w))
188         delta = np.dot(X.T, y_hat - y) / N
189         return delta
190
191 class KFoldValidation:
192     def __init__(self, folds, path, lr, max_iter, epsilon, beta):
193         self.folds = folds
194         self.data = pd.read_csv(path)
195         self.lr = lr
196         self.max_iter = max_iter
197         self.epsilon = epsilon
198         self.beta = beta
199
200     def k_fold_validation(self, normalize=False, inc_od=False, order=3):
201         '''
202             This function performs the k-fold validation
203
204             normalize - whether to perform normalization
205             inc_od - whether to increase the feature order
206             order - the order of the added feature
207         '''
208         folds = self.folds
209         data = self.data
210         accuracies = []
211
212         if inc_od:
213             data = self.rise_order(data, order)
214
215         log_reg = LogisticRegression(data=data, folds=self.folds, lr=self.lr,
max_iter=self.max_iter, beta=self.beta, epsilon=self.epsilon)
216
217         log_reg.shuffle_data()
218
219         for fold in range(folds):
220             t_X, t_y, v_X, v_y = log_reg.partition(fold)
221             # t_X --> test value X, v_X --> validation value X
222             gradient_val, t_acc_val, v_acc_val, best_w = log_reg.fit(t_X, t_y
, v_X, v_y, normalize=normalize)
223
224             accuracies.append(np.max(v_acc_val))
225
226             # Uncomment this block to display the accuracy diagram
227             plt.figure()
228             plt.plot(t_acc_val, label = 'Training accuracy')
229             plt.plot(v_acc_val, label='Validation accuracy')
230             plt.axvline(log_reg.marker, color='r', label='Best Weights')
231             plt.xlabel('Iteration Number')
232             plt.ylabel('Accuracy')
233             plt.legend()
234             plt.show()
235             print("Learning Rate: " + str(log_reg.lr))
236             print("Average Accuracy: "+str(np.mean(accuracies)))
237
238             # Uncomment this block to display the gradient diagram
239             plt.figure()
240             plt.plot(gradient_val)

```

```

241         plt.xlabel('Iteration Number')
242         plt.ylabel('Gradient')
243         plt.show()
244         print("-----")
245
246         mean_acc = np.mean(accuracies)
247         return mean_acc
248
249     def rise_order(self, data, order=3):
250         ret_val = data
251         for i in range(2, order + 1):
252             data_powered = data.pow(i)
253             ret_val = ret_val.iloc[:, :-1]
254             ret_val = pd.concat([ret_val, data_powered], axis=1)
255         return ret_val
256
257     """### Perform 10-fold validation for Hepatitis dataset"""
258
259     path = "/content/drive/My Drive/ECSE_551_F_2020/Mini_Project_01/hepatitis.csv"
260     dataset_name = "Hepatitis"
261     default_lr = 0.01
262     default_max_iter = 10000
263     default_epsilon = 5e-3
264     default_beta = 0.99
265
266     # the input is the optimum hyperparameters found during testing
267     hepatitis_learning = KFoldValidation(folds=10, path=path, lr=default_lr,
268                                         max_iter=default_max_iter, epsilon=default_epsilon, beta=default_beta)
269     # the input is the optimum model found during testing
270     mean_acc = hepatitis_learning.k_fold_validation()
271
272     """### Perform 10-fold validation for Bankruptcy dataset"""
273
274     path = "/content/drive/My Drive/ECSE_551_F_2020/Mini_Project_01/bankruptcy.csv"
275     dataset_name = "Bankruptcy"
276     default_lr = 0.1
277     default_max_iter = 25000
278     default_epsilon = 1e-3
279     default_beta = 0.99
280
281     # the input is the optimum hyperparameters found during testing
282     bankruptcy_learning = KFoldValidation(folds=10, path=path, lr=default_lr,
283                                           max_iter=default_max_iter, epsilon=default_epsilon, beta=default_beta)
284     # the input is the optimum model found during testing
285     mean_acc = bankruptcy_learning.k_fold_validation(normalize=True, inc_od=True)

```