
ECSE 551 Fall 2020 Mini-project 3 Report

Fei Peng
260712440
fei.peng@mail.mcgill.ca

Yukai Zhang
260710915
yukai.zhang@mail.mcgill.ca

Yi Zhu
260716006
yi.zhu6@mail.mcgill.ca

Abstract

Image identification and classification are two of the most popular research areas in machine learning as well as computer vision. In this mini-project, we chose to investigate and evaluate the performance of different Convolutional Neural Network (i.e., CNN) models on a large image dataset. To have a better understanding of the structure of CNN, we implemented a model from scratch and analyzed its behavior. In addition, we tested different CNN layer arrangements and examined the effect of hyperparameters (e.g., batch size, number of epochs, and learning rate) before the optimum combination was selected, which yields the highest accuracy in the validation set. The performance of other CNN implementations provided by PyTorch and TensorFlow were also investigated, from which we found that VGG has the highest accuracy without any pre-trained data. According to our testing result, correctly choosing a proper CNN layer arrangement can greatly boost the result accuracy, avoid allocating too much memory as well as make sure the training process converges after reasonable iterations. Details will be presented in the following parts of this report.

1 Introduction

The increasing popularity of computer vision and the advance of machine learning technique have brought the opportunities to perform apparel classification [1]. The main objective of this mini-project is to apply machine learning methods on classification of clothes from five different categories where each has a distinctive price from \$1 to \$5. There are 60,000 monochrome images provided in the training dataset for this mini-project, and each image consists of three clothes. Therefore, their total price lies between \$5 and \$13 which has 9 different values. After applying machine learning methods on the training dataset, an additional 10,000 images are used to test the accuracy of the classification techniques and the results are submitted on Kaggle competition.

In this project, we are allowed to freely choose any machine learning methods, so as to explore their performance. However, the three clothes could appear in any random place in each image, and most of the classification methods we have been using so far in this course are sensitive to this position shift. Thus, a fully connected layer analysis (e.g., fully connected neural networks) is preferred as it can be used to achieve shift invariance, but performing this method requires a large number of training instances (i.e., weight parameters) as well as a long learning period, whereas we have limited computational resources and we are forbidden to import any pre-trained data [2]. Convolutional Neural Network, on the other hand, conserves the benefit of shift invariant by applying the same weight configuration across the image space. A convolution neural network consists of one or multiple convolutional layers, which are followed by fully connected layers. This network arrangement allows CNN to process/analyze an image in a similar manner as the biological neural systems.

Since the CNN method is wildly applied to solve many image-recognition problems, different types of network designs are available. Choosing the one that best fits this mini-project is essential. Thus, we investigated the performance of the CNN we implemented ourselves, and ResNet18 [3], AlexNet [4], VGG16 as well as VGG19 (both versions provided by PyTorch and TensorFlow) [5].

After comparing the validation accuracy of each implementation of CNN, VGG19 (by TensorFlow) was found to show a promising result. It yielded an accuracy of 99% in the validation dataset and 97% in the Kaggle competition.

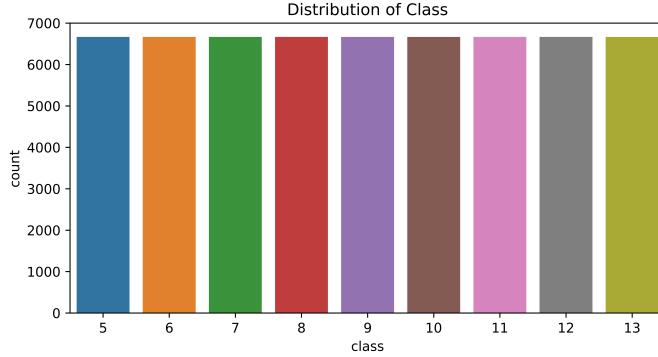


Figure 1: Class distribution of train data

2 Datasets

Before performing any investigation on the supervised machine learning and image classification models, it is important that the characteristics of the dataset (i.e., both features and classes) are properly analyzed and understood. Aside from that, data preprocessing lays the groundwork for the following data analysis, as its product is the final training set [6].

The training dataset we used is in *Train.pkl* which contains 60,000 grayscale image samples, and their corresponding classes are imported from *TrainLabels.csv*. The test dataset is 10,000 grayscale image samples contained in *Test.pkl*, from which we predicted the classes and submitted the result to Kaggle competition.

In this project, the following data analyzation and preprocessing had been conducted.

2.1 Feature Analysis

The dataset is a modified version of the Fashion-MNIST dataset, which consists of samples with 641281 pixel values as features representing grayscale images. Therefore, the total number of features in each sample is 8192. As the number of samples in this dataset is numerous (i.e., 60,000 training samples), we considered using deep learning to achieve optimum accuracy with large dataset [7], from which we chose Convolutional Neural Networks (i.e., CNN) instead of fully connected neural networks for fewer memory requirements and lower latency [2]. Moreover, based on our previous experience, the images are relatively large and complex, therefore we anticipated to use deeper neural networks (i.e., more layers of neurons) for better accuracy. Details will be discussed later in Section 3.

2.2 Target Analysis

In this dataset, each image contains three articles, and our goal is to predict the total price of all articles presented in the image. The prices (i.e., classes) range from 5 to 13, therefore, it is a nine-way classification problem. The distribution of each class is shown in Figure 1. The uncertainty in prediction can be quantified by the entropy of the dataset with Equation 1, and the calculated data entropy is 3.17 which stands for an even distribution. Therefore, the number of samples in each class is equal. Moreover, before feeding the dataset to any training process, we need to make sure that the class values are encoded. Since the prices are integer numbers range from 5 to 13, we could simply encode them by subtracting the values by 5, thus the obtained target values would range from 0 to 8.

$$H(X) = - \sum_{x \in X} p(x) \log_2 p(x) \quad (1)$$

2.3 Normalization

The original pixel value of the images in the dataset is between 0 and 255, to centralize the value distribution while keeping its precision, we normalize the data with mean and standard division both equal to 0.5. Therefore, after normalization, the pixel values range from -1 to 1, which removes any scale factors that might influence the classification result [8].

3 Proposed Approach

In this section, we are going to discuss the models we implemented and the methods we used for this project, as well as the model and method selection in terms of accuracy, robustness/consistency, and speed. However, due to time constraint of this project, the reported accuracy and time might not be based on the absolute optimum hyperparameters, but rather try and error.

3.1 Training/Validation Splits

As mentioned previously in Section 2, the training dataset contains 60,000 samples, while the testing dataset contains 10,000 samples, however, only the training dataset contains the corresponding target classes. Therefore, for model selection, we chose to split the training dataset to 50,000 training samples and 10,000 testing samples, in order to perform validation. In this way, we could ensure consistency between the validation and testing process, and increase the possibility to reach higher accuracy on predicting the testing dataset. Due to the fact that deep learning methods are relatively more time consuming than normal machine learning algorithms, k-fold validation technique is not as common, thus, we compared the accuracy of different models based only on those 10,000 testing samples that are separated from the training dataset.

3.2 Hyperparameters

It is common practice to decay the learning rate in order to achieve better accuracy. However, one can usually obtain the same learning curve on both training and test sets by instead increasing the batch size during training [9]. As will be explained detailed later in Section 4, a smaller learning rate leads to a longer training time to approach the same accuracy. In CNN, it means that more epochs are required for training the network to preserve the same accuracy. Since the time consumed in each epoch is relatively constant, there is a linear relationship between number of epochs and total training time. However, increasing the learning rate reduces the number of parameter updates, thus, leads to greater parallelism and shorter training times. To compensate the influence of increasing the learning rate (e.g., by a factor of ϵ), we could increase the batch size by a factor (e.g., B) that is linearly proportional to the increase of learning rate (i.e., $B \propto \epsilon$) [9]. However, increasing the batch size alone could not ensure the same accuracy [10], whereas in our case decrease the accuracy as will be explained later in Section 4. Therefore, by tuning the learning rate and batch size together, we could select the most suitable combination to achieve higher accuracy and shorter training time. Moreover, we could increase the momentum coefficient (i.e., β) in gradient descent and scale $B \propto \frac{1}{1-\beta}$, while maintaining the same accuracy [9].

On the other hand, although increasing the number of epochs leads to longer training time, it also helps with achieving higher accuracy. In real life, it is usually a trade-off between choosing shorter training time and higher accuracy. In our case, since this is a competition on test accuracy, we would choose higher accuracy by sacrificing training time. However, accuracy will eventually converge, and continuing to train the model for more epochs will cause overtraining.

3.3 Algorithm Selection

At the beginning of this project, we started with building a Convolutional Neural Network (i.e., CNN) from scratch, in order to explore the structure and characteristics of the network. At the beginning, we configured a rather simple network, which starts with only four 2D convolution layers, each followed by a ReLU activation function as well as a 2D max pooling, and ends with two fully connected layers. However, we underestimated the requirement of the feature size (i.e., number of pixels in each image) on the depth of the network. The accuracy did not advance with the increasing number of epochs, but remained no difference from random classification, and could easily end up with data underflow/overflow, no matter how we changed the gradient descent and image normalization parameters. Therefore, we decided to increase the depth of our CNN model by learning and imitating one of the existing successful models provided by PyTorch [11], which is ResNet18 (i.e., residual network) [3], and tweaked the structure to suit our use. As a result, our CNN model contains twenty 2D convolution layers, followed by batch normalisation and ReLU activation functions, as well as one fully connected layer at the end. With this model, the running loss (i.e., cross entropy loss) started to converge as the number of epochs increased, and the accuracy began to rise.

Next, we continued to investigate the behavior of CNN models by performing experiments on ResNet18 [3], AlexNet [4], VGG16, VGG19 as well as VGG19 with batch normalization (i.e., VGG-19 BN) [5] provided by PyTorch [11], as well as VGG19 from TensorFlow (i.e., VGG-19 TF) [12], with minor adjustment on the models to fit our situation (i.e., input color channel equals one, and number of output class equals nine). AlexNet is one of the first Deep Convolutional Networks to achieve considerable accuracy, which consists of five convolutional layers and three fully connected layers, with ReLU as the activation function. VGG16 and VGG19 are both VGGNet, which differ only in the total number of layers in the network (i.e., 16 and 19 layers respectively). The idea of VGG is to use fixed size kernels so that all the variable size convolutional kernels used in AlexNet can be replicated by making use of multiple 3x3 kernels as building blocks. As a result, VGG reduces the number of parameters and the training time. ResNet18 are 18-layer residual nets, which have shortcut connection added to each pair of 3x3 filters compared to plain nets that are inspired by VGGNet. Thus, ResNet is able to find simpler mappings when they exist. Later in the next section (i.e., Section 4), we are going to discuss in detail the experiment results.

4 Results

When performing machine learning with CNN, multiple epochs are needed in order to obtain the highest accuracy. After evaluating the time and accuracy with respect to the number of epochs for the CNN implemented by ourselves, we found that the time required for each epoch is approximately constant and the total time is linearly proportional to the number of epochs. Accuracy, on the other hand, has an abrupt increase at the first several epochs. The slope of the curve gradually reduces as the accuracy approaches 1.0. Details are shown in Figure 2. After training for 32 epochs, our model achieved an accuracy of 84% and the training time is 995.74 seconds, which are shown in Table 1.

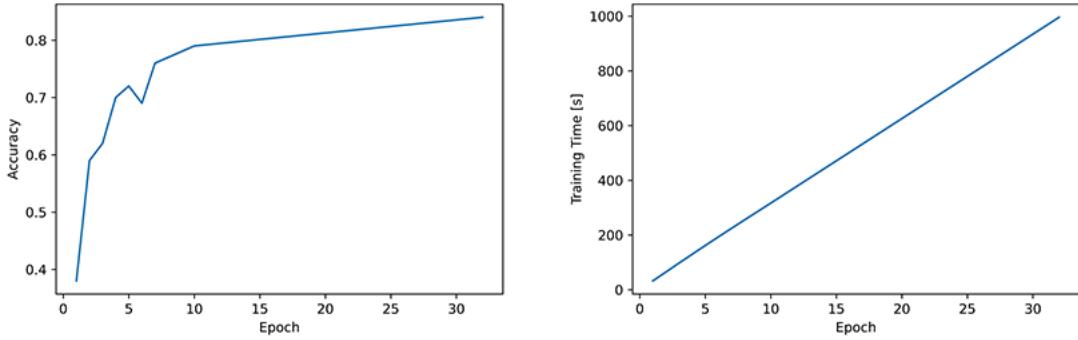


Figure 2: Accuracy (left) and training time (right) with respect to number of epochs

Table 1: Accuracy and training time for seven CNN models (32 epochs)

	CNN	ResNet-18	AlexNet	VGG-16	VGG-19	VGG-19 BN	VGG-19 TF
Accuracy	84%	87%	84%	94%	98%	98%	99%
Time (s)	995.74	1598.56	1275.41	5399.81	6151.02	6758.24	6085.63

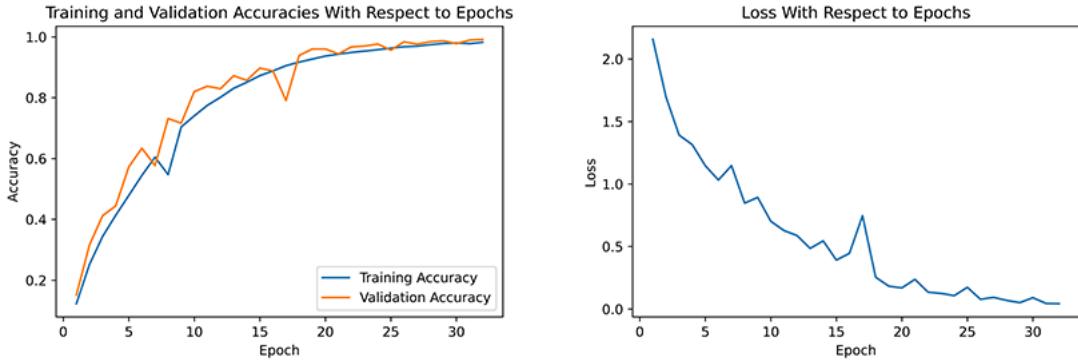


Figure 3: Training/validation accuracy (left) and loss (right) with respect to number of epochs

We then conducted experiments with other CNN models. The performances are recorded in Table I where 32 epochs of training are executed for each model. Despite having a slightly lower accuracy (84%), the CNN implemented by ourselves has demonstrated an incredible efficiency which takes only 15% of the training time compared to more complex models such as VGG. The VGG19 provided by TensorFlow, on the other hand, can yield the highest accuracy (99%) with a proper combination of hyperparameters. In our case, we chose SGD as our gradient descent method with a learning rate of 0.001 and a momentum of 0.7. We also set the loss function to *sparse_categorical_crossentropy* and the metrics we use to evaluate the model is *accuracy*. We disabled the usage of pre-train data by setting the weight of VGG19 to be *None* and we specified a batch size of 16. Its result was used in Kaggle competition and achieved a leaderboard accuracy of 0.98366.

To have a better understanding of this model (i.e., VGG-19 TF), we did further investigation. We picked 10,000 of the 60,000 training dataset as the validation set and the relationship between training/validation accuracies as well as loss with respect to epochs are shown in Figure 3. It is noticeable that the curve of validation accuracy keeps rising in the last few epochs and we can infer that overtraining has not occurred within the first 32 epochs. The loss decreases and approaches 0 as the number of epochs increases.

The accuracies and convergence time depend on not only the training models, but also the hyperparameters. The effects of the learning rate/momentum as well as batches size are investigated. We used VGG19 and VGG19 with batch normalization to test the effect of learning rate and momentum. Their accuracies after 32 epochs are shown in Table 2 and their loss curves are plotted in Figure 4. It appears that with a higher learning rate, the loss will have a more rapid drop in the first several epochs. The loss difference between different learning rates will decrease as the number of epochs gets larger.

Table 2: Accuracy performance of two CNN models under different hyperparameter sets (32 epochs)

	VGG-19	VGG-19 BN
Learning Rate = 0.001 & Momentum = 0	56%	68%
Learning Rate = 0.005 & Momentum = 0.5	98%	98%

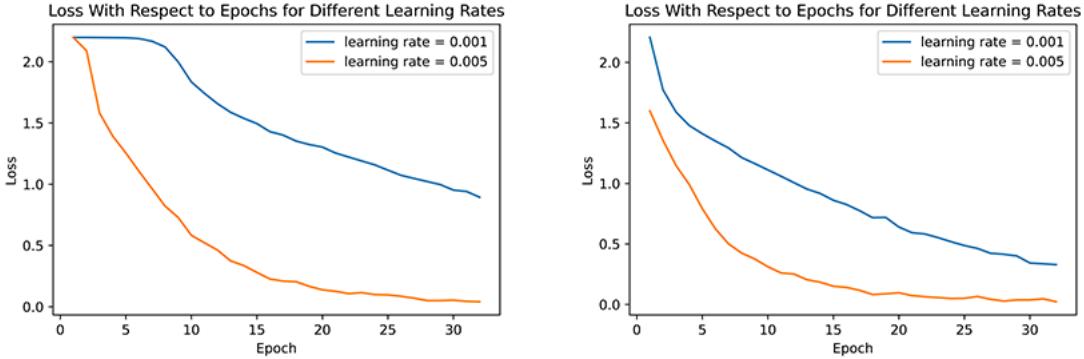


Figure 4: Loss with respect to epochs in different learning rates in VGG19 (left) and VGG19 BN (right)

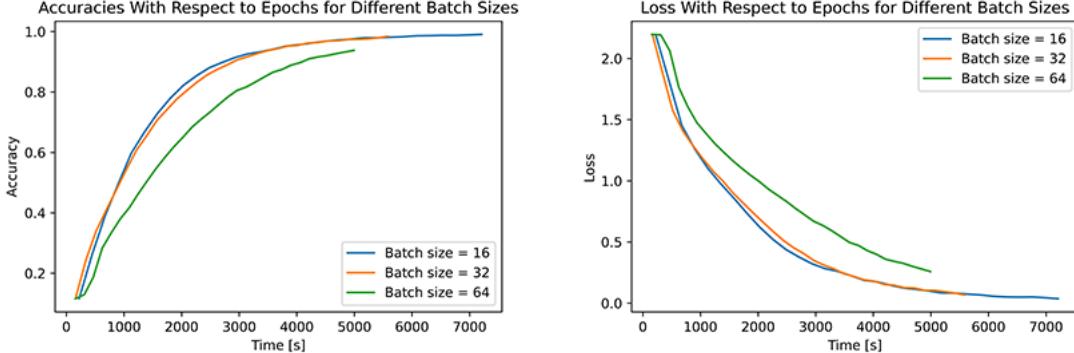


Figure 5: Accuracy (left) and loss (right) with respect to number of epochs for different batch sizes

5 Discussion and Conclusion

In conclusion, the Convolutional Neural Network model we implemented from scratch in this project performed as expected for classifying the clothes images. It achieved a validation accuracy of 84% after training for 32 epochs and the training process takes 995.74 seconds. However, the CNN models provided by TensorFlow yields the highest validation accuracy (99%) without using any pretrained data. The validation accuracy is found to rise at a decreasing rate as the number of epochs grows. There does not exist a sign of overtraining within the first 32 epochs when the learning rate is set to be no larger than 0.005.

The future investigation could be to improve the layer arrangement of our self-implemented CNN and to find the best combination of hyperparameters in order to achieve a higher accuracy with less training time/epochs.

On the other hand, we achieved a leaderboard accuracy of 0.98366 using the VGG19 provided by TensorFlow without pre-trained data. 64 epochs of training were performed with a batch size of 16 in order to achieve this result as the validation accuracy barely changes (i.e., converges) at this point. In the future, we will take a closer look at the hyperparameters of the VGG19 algorithm to find the best combination that suits this machine learning problem.

6 Statement of Contributions

Yukai Zhang was in charge of researching the structure of the Convolutional Neural Network and implementing a CNN from scratch. Fei Peng browsed the models in PyTorch and TensorFlow, modified and trained the selected base models, as well as investigated the relationship between accuracy and number of epochs. Yi Zhu was responsible for analyzing the correlation between time and number of epochs, as well as proofreading the report.

References

- [1] L. Bossard, M. Dantone, C. Leistner, C. Wengert, T. Quack, and L. Van Gool, "Apparel Classification with Style," Berlin, Heidelberg, 2013: Springer Berlin Heidelberg, in Computer Vision – ACCV 2012, pp. 321-335.
- [2] S. Hijazi, R. Kumar, and C. Rowen, "Using Convolutional Neural Networks for Image Recognition By," 2015.
- [3] K. He, X. Zhang, S. Ren and J. Sun, "Deep Residual Learning for Image Recognition," *2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, Las Vegas, NV, 2016, pp. 770-778, doi: 10.1109/CVPR.2016.90.
- [4] A. Krizhevsky, "One weird trick for parallelizing convolutional neural networks," *ArXiv*, vol. abs/1404.5997, 2014.
- [5] K. Simonyan and A. Zisserman, "Very Deep Convolutional Networks for Large-Scale Image Recognition," *CoRR*, vol. abs/1409.1556, 2015.
- [6] S. Kotsiantis, D. Kanellopoulos, and P. Pintelas, "Data Preprocessing for Supervised Learning," *World Academy of Science, Engineering and Technology, International Journal of Computer, Electrical, Automation, Control and Information Engineering*, vol. 1, pp. 4104-4109, 2007.
- [7] Y. LeCun, Y. Bengio, and G. Hinton, "Deep learning," *Nature*, vol. 521, no. 7553, pp. 436-444, 2015/05/01 2015, doi: 10.1038/nature14539.
- [8] G. Finlayson, B. Schiele, and J. Crowley, "Comprehensive Colour Image Normalization," in *ECCV*, 1998.
- [9] S. L. Smith, P. Kindermans, and Q. V. Le, "Don't Decay the Learning Rate, Increase the Batch Size," *ArXiv*, vol. abs/1711.00489, 2018.
- [10] Y. You, I. Gitman, and B. Ginsburg, "Scaling SGD Batch Size to 32K for ImageNet Training," *ArXiv*, vol. abs/1708.03888, 2017.
- [11] A. Paszke et al., "PyTorch: An Imperative Style, High-Performance Deep Learning Library," *ArXiv*, vol. abs/1912.01703, 2019.
- [12] M. Abadi et al., "TensorFlow: Large-Scale Machine Learning on Heterogeneous Distributed Systems," *ArXiv*, vol. abs/1603.04467, 2016.

Appendix

(Please see following pages attached for code implementations)

```

1 # -*- coding: utf-8 -*-
2 """Data_Preprocessing.ipynb
3
4 Automatically generated by Colaboratory.
5
6 Original file is located at
7     https://colab.research.google.com/drive/1ALc0_L0DWWHcrm6Qcxv-KH4TzpSpAlak
8
9 <center><h1>Mini Project 3 - Convolutional Neural Network</h1>
10 <h3>Data Preprocessing</h3>
11 <h4>This file performs some of the operations on Data Preprocessing and
Analysis.</h4></center>
12
13 <h3>Team Members:</h3>
14 <center>
15 Yi Zhu, 260716006<br>
16 Fei Peng, 260712440<br>
17 Yukai Zhang, 260710915
18 </center>
19
20 # Importations
21 """
22
23 from google.colab import drive
24 drive.mount('/content/drive')
25
26 # make path = './' in-case you are running this locally
27 path = '/content/drive/My Drive/ECSE_551_F_2020/Mini_Project_03/'
28
29 import numpy as np
30 import pandas as pd
31 import matplotlib.pyplot as plt
32 import seaborn as sns
33 from scipy import stats
34 from google.colab import files
35 from sklearn.preprocessing import LabelEncoder
36 from scipy.stats import entropy
37
38 """# Data Preprocessing"""
39
40 dataset = pd.read_csv(path+"TrainLabels.csv")
41 # reddit_test = pd.read_csv(path+"test.csv")
42
43 y = dataset['class']
44
45 class Data_Processing:
46     def __init__(self, data, name='New Data'):
47         self.data = data
48         self.name = name
49
50     def show_y_dist(self, ydata):
51         plt.figure(figsize=(8,4))
52         plt.subplot(111), sns.countplot(x='class', data=ydata)
53         plt.title('Distribution of Class')
54         plt.savefig("Distribution of Class.png", dpi=1200)
55         files.download("Distribution of Class.png")
56         plt.show()
57
58 data_analysis = Data_Processing(dataset.values, 'TrainLabels.csv')
59 data_analysis.show_y_dist(dataset)
60
61 # calculate the data entropy
62 le = LabelEncoder() # encoder for classes

```

```
63 le.fit(y)
64 y_label = le.transform(y)
65 n_k = len(le.classes_)
66 N = len(y)
67 theta_k = np.zeros(n_k) # probability of class k
68 # compute theta values
69 for k in range(n_k):
70     count_k = (y_label==k).sum()
71     theta_k[k] = count_k / N
72
73 print("Data entropy is", entropy(theta_k, base=2))
```

File - E:\Study\ECSE551\Mini_Project_3\Code\cnn.py

```

57 # Dataloader class
58 Although we can access all the training data using the Dataset class, for
      neural networks, we would need batching, shuffling, multiprocess data loading
      , etc. DataLoader class helps us to do this. The DataLoader class accepts a
      dataset and other parameters such as batch_size.
59 """
60
61 # Transforms are common image transformations. They can be chained together
      using Compose.
62 # Here we normalize images img=(img-0.5)/0.5
63 img_transform = transforms.Compose([
64     transforms.ToTensor(),
65     transforms.Normalize((0.5,), (0.5,))
66 ])
67
68 # img_file: the pickle file containing the images
69 # label_file: the .csv file containing the labels
70 # transform: We use it for normalizing images (see above)
71 # idx: This is a binary vector that is useful for creating training and
      validation set.
72 # It return only samples where idx is True
73
74 class MyDataset(Dataset):
75     def __init__(self, img_file, label_file, transform=None, idx = None):
76         self.data = pickle.load(open( img_file, 'rb' ), encoding='bytes')
77         self.targets = np.genfromtxt(label_file, delimiter=',', skip_header=1
    )[:,1:]
78         if idx is not None:
79             self.targets = self.targets[idx]
80             self.data = self.data[idx]
81             self.transform = transform
82             self.targets -= 5
83
84     def __len__(self):
85         return len(self.targets)
86
87     def __getitem__(self, index):
88         img, target = self.data[index], int(self.targets[index])
89         img = Image.fromarray(img.astype('uint8'), mode='L')
90
91         if self.transform is not None:
92             img = self.transform(img)
93
94         return img, target
95
96 # Read image data and their label into a Dataset class
97 dataset = MyDataset('./Train.pkl', './TrainLabels.csv', transform=
    img_transform, idx=None)
98
99 batch_size = 32 #feel free to change it
100 dataloader = DataLoader(dataset, batch_size=batch_size, shuffle=True)
101
102 # Read a batch of data and their labels and display them
103 # Note that since data are transformed, they are between [-1,1]
104 imgs, labels = (next(iter(dataloader)))
105 imgs = np.squeeze(imgs)
106 plt.imshow(imgs[5].cpu().numpy(), cmap='gray', vmin=-1, vmax=1) #.transpose()
107
108 """# CNN"""
109
110 import torch.nn as nn
111 import torch.nn.functional as F
112

```

```

113 import torch.optim as optim
114
115 # cnn
116 # This cnn is based on the structure of resnet18
117 class Net(nn.Module):
118     def __init__(self):
119         super(Net, self).__init__()
120         self.conv1 = nn.Conv2d(1, 32, kernel_size=(7,7), stride=(2, 2),
121                             padding=(3, 3), bias=False)
122         self.bn1 = nn.BatchNorm2d(32, eps=1e-5, momentum=0.1, affine=True,
123                                track_running_stats=True)
124         self.relu = nn.ReLU(inplace=True)
125         self.maxpool = nn.MaxPool2d(kernel_size=3, stride=2, padding=1,
126                                    dilation=1, ceil_mode=False)
127
128         # Layer 1
129         self.layer1block1conv1 = nn.Conv2d(32, 32, kernel_size=(3, 3), stride
130 = (1, 1), padding=(1, 1), bias=False)
131         self.layer1block1bn1 = nn.BatchNorm2d(32, eps=1e-05, momentum=0.1,
132                                affine=True, track_running_stats=True)
133         self.layer1block1relu = nn.ReLU(inplace=True)
134         self.layer1block1conv2 = nn.Conv2d(32, 32, kernel_size=(3, 3), stride
135 = (1, 1), padding=(1, 1), bias=False)
136         self.layer1block1bn2 = nn.BatchNorm2d(32, eps=1e-05, momentum=0.1,
137                                affine=True, track_running_stats=True)
138
139         # Layer 2
140         self.layer2block1conv1 = nn.Conv2d(32, 64, kernel_size=(3, 3), stride
141 = (2, 2), padding=(1, 1), bias=False)
142         self.layer2block1bn1 = nn.BatchNorm2d(64, eps=1e-05, momentum=0.1,
143                                affine=True, track_running_stats=True)
144         self.layer2block1relu = nn.ReLU(inplace=True)
145         self.layer2block1conv2 = nn.Conv2d(64, 64, kernel_size=(3, 3), stride
146 = (1, 1), padding=(1, 1), bias=False)
147         self.layer2block1bn2 = nn.BatchNorm2d(64, eps=1e-05, momentum=0.1,
148                                affine=True, track_running_stats=True)
149         self.layer2block1conv3 = nn.Conv2d(64, 64, kernel_size=(1, 1), stride
150 = (2, 2), bias=False)
151         self.layer2block1bn3 = nn.BatchNorm2d(64, eps=1e-05, momentum=0.1,
152                                affine=True, track_running_stats=True)
153
154         # Layer 3
155         self.layer3block1conv1 = nn.Conv2d(64, 128, kernel_size=(3, 3), stride

```

```

154 =(2, 2), padding=(1, 1), bias=False)
155         self.layer3block1bn1 = nn.BatchNorm2d(128, eps=1e-05, momentum=0.1,
156         affine=True, track_running_stats=True)
157         self.layer3block1relu = nn.ReLU(inplace=True)
158         self.layer3block1conv2 = nn.Conv2d(128, 128, kernel_size=(3, 3),
159         stride=(1, 1), padding=(1, 1), bias=False)
160         self.layer3block1bn2 = nn.BatchNorm2d(128, eps=1e-05, momentum=0.1,
161         affine=True, track_running_stats=True)
162         self.layer3block1conv3 = nn.Conv2d(128, 128, kernel_size=(1, 1),
163         stride=(2, 2), bias=False)
164         self.layer3block1bn3 = nn.BatchNorm2d(128, eps=1e-05, momentum=0.1,
165         affine=True, track_running_stats=True)
166
167         self.layer3block2conv1 = nn.Conv2d(128, 128, kernel_size=(3, 3),
168         stride=(1, 1), padding=(1, 1), bias=False)
169         self.layer3block2bn1 = nn.BatchNorm2d(128, eps=1e-05, momentum=0.1,
170         affine=True, track_running_stats=True)
171         self.layer3block2relu = nn.ReLU(inplace=True)
172         self.layer3block2conv2 = nn.Conv2d(128, 128, kernel_size=(3, 3),
173         stride=(1, 1), padding=(1, 1), bias=False)
174         self.layer3block2bn2 = nn.BatchNorm2d(128, eps=1e-05, momentum=0.1,
175         affine=True, track_running_stats=True)
176
177         # Layer 4
178         self.layer4block1conv1 = nn.Conv2d(128, 256, kernel_size=(3, 3),
179         stride=(2, 2), padding=(1, 1), bias=False)
180         self.layer4block1bn1 = nn.BatchNorm2d(256, eps=1e-05, momentum=0.1,
181         affine=True, track_running_stats=True)
182         self.layer4block1relu = nn.ReLU(inplace=True)
183         self.layer4block1conv2 = nn.Conv2d(256, 256, kernel_size=(3, 3),
184         stride=(1, 1), padding=(1, 1), bias=False)
185         self.layer4block1bn2 = nn.BatchNorm2d(256, eps=1e-05, momentum=0.1,
186         affine=True, track_running_stats=True)
187
188         self.layer4block2conv1 = nn.Conv2d(256, 256, kernel_size=(3, 3),
189         stride=(1, 1), padding=(1, 1), bias=False)
190         self.layer4block2bn1 = nn.BatchNorm2d(256, eps=1e-05, momentum=0.1,
191         affine=True, track_running_stats=True)
192
193         self.avgpool = nn.AdaptiveAvgPool2d(output_size=(1,1))
194         self.fc = nn.Linear(in_features=256, out_features=9, bias=True)
195
196         def forward(self, x):
197             x = self.conv1(x)
198             x = self.bn1(x)
199             x = self.relu(x)
200             x = self.maxpool(x)
201
202             # Layer 1
203             x = self.layer1block1conv1(x)
204             x = self.layer1block1bn1(x)
205             x = self.layer1block1relu (x)
206             x = self.layer1block1conv2(x)
207             x = self.layer1block1bn2(x)

```

```

198
199         x = self.layer1block2conv1(x)
200         x = self.layer1block2bn1(x)
201         x = self.layer1block2relu (x)
202         x = self.layer1block2conv2(x)
203         x = self.layer1block2bn2(x)
204
205     # Layer 2
206     x = self.layer2block1conv1(x)
207     x = self.layer2block1bn1(x)
208     x = self.layer2block1relu(x)
209     x = self.layer2block1conv2(x)
210     x = self.layer2block1bn2(x)
211     x = self.layer2block1conv3(x)
212     x = self.layer2block1bn3(x)
213
214     x = self.layer2block2conv1(x)
215     x = self.layer2block2bn1(x)
216     x = self.layer2block2relu(x)
217     x = self.layer2block2conv2(x)
218     x = self.layer2block2bn2(x)
219
220     # Layer 3
221     x = self.layer3block1conv1(x)
222     x = self.layer3block1bn1(x)
223     x = self.layer3block1relu(x)
224     x = self.layer3block1conv2(x)
225     x = self.layer3block1bn2(x)
226     x = self.layer3block1conv3(x)
227     x = self.layer3block1bn3(x)
228
229     x = self.layer3block2conv1(x)
230     x = self.layer3block2bn1(x)
231     x = self.layer3block2relu(x)
232     x = self.layer3block2conv2(x)
233     x = self.layer3block2bn2(x)
234
235     # Layer 4
236     x = self.layer4block1conv1(x)
237     x = self.layer4block1bn1(x)
238     x = self.layer4block1relu(x)
239     x = self.layer4block1conv2(x)
240     x = self.layer4block1bn2(x)
241     x = self.layer4block1conv3(x)
242     x = self.layer4block1bn3(x)
243
244     x = self.layer4block2conv1(x)
245     x = self.layer4block2bn1(x)
246     x = self.layer4block2relu(x)
247     x = self.layer4block2conv2(x)
248     x = self.layer4block2bn2(x)
249
250     x = self.avgpool(x)
251     x = torch.flatten(x, 1)
252     x = self.fc(x)
253
254     return x
255
256 train_index = np.arange(50000)
257 test_index = np.arange(50000, 60000)
258 batch_size = 32 #feel free to change it
259
260 # Read image data and their label into a Dataset class

```

```

261 train_set = MyDataset('./Train.pkl', './TrainLabels.csv', transform=
262     img_transform, idx=train_index)
263 test_set = MyDataset('./Train.pkl', './TrainLabels.csv', transform=
264     img_transform, idx=test_index)
265 train_loader = DataLoader(train_set, batch_size=batch_size, shuffle=True,
266     num_workers=2)
267 test_loader = DataLoader(test_set, batch_size=batch_size, shuffle=True,
268     num_workers=2)
269
270 net = Net()
271 # if there is a available cuda device, use GPU, else, use CPU
272 device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
273 net = net.to(device)
274
275 # set criterion to cross entropy loss
276 criterion = nn.CrossEntropyLoss()
277 # set learning rate to 0.001
278 optimizer = optim.SGD(net.parameters(), lr=0.01, momentum=0.5)
279
280 running_loss = 0.0
281 num_epochs = 32
282 for epoch in range(num_epochs):
283     for i, data in enumerate(train_loader, 0):
284         img, label = data
285         img = img.to(device)
286         label = label.to(device)
287
288         # zero the parameter gradients
289         optimizer.zero_grad()
290
291         # forward + backward + optimize
292         outputs = net(img)
293         loss = criterion(outputs, label)
294         loss.backward()
295         optimizer.step()
296
297         # print statistics
298         running_loss += loss.item()
299         if i % 320 == 319: # print every 320 mini-batches
300             print('[%d, %5d] loss: %.3f' %
301                 (epoch + 1, i + 1, running_loss / 320))
302             running_loss = 0.0
303
304             torch.save(net.state_dict(), '/model.pth')
305             torch.save(optimizer.state_dict(), '/optimizer.pth')
306
307 print('Finished Training')
308
309 correct = 0
310 total = 0
311
312 # calculate accuracy
313 with torch.no_grad():
314     for data in test_loader:
315         images, labels = data
316         images = images.to(device)
317         labels = labels.to(device)
318         outputs = net(images)
319         # get the index of the max output
320         _, predicted = torch.max(outputs.data, 1)
321         total += labels.size(0)
322         correct += (predicted == labels).sum().item()
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593
594
595
596
597
598
599
600
601
602
603
604
605
606
607
608
609
610
611
612
613
614
615
616
617
618
619
620
621
622
623
624
625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701
702
703
704
705
706
707
708
709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
781
782
783
784
785
786
787
788
789
790
791
792
793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809
810
811
812
813
814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863
864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
888
889
890
891
892
893
894
895
896
897
898
899
900
901
902
903
904
905
906
907
908
909
910
911
912
913
914
915
916
917
918
919
920
921
922
923
924
925
926
927
928
929
930
931
932
933
934
935
936
937
938
939
940
941
942
943
944
945
946
947
948
949
950
951
952
953
954
955
956
957
958
959
960
961
962
963
964
965
966
967
968
969
970
971
972
973
974
975
976
977
978
979
980
981
982
983
984
985
986
987
988
989
990
991
992
993
994
995
996
997
998
999
999

```

```
320 print('Accuracy of the network: %d %%' % (
321     100 * correct/ total))
```

```

1 # -*- coding: utf-8 -*-
2 """PyTorch_Net.ipynb
3
4 Automatically generated by Colaboratory.
5
6 Original file is located at
7     https://colab.research.google.com/drive/1VaRJXCTRYPnSrfmQcq6PVJuMZ5xxB250
8
9 <center><h1>Mini Project 3 - Convolutional Neural Network</h1>
10 <h4>The PyTorch File.</h4></center>
11
12 <h3>Team Members:</h3>
13 <center>
14 Yi Zhu, 260716006<br>
15 Fei Peng, 260712440<br>
16 Yukai Zhang, 260710915
17 </center>
18 """
19
20 # Commented out IPython magic to ensure Python compatibility.
21 from google.colab import drive
22 drive.mount("/content/drive")
23
24 # %cd '/content/drive/My Drive/ECSE_551_F_2020/Mini_Project_03/'
25
26 import numpy as np
27 import pandas as pd
28 import torch
29 import torchvision
30 import torchvision.transforms as transforms
31 import matplotlib.pyplot as plt
32 import pickle
33
34 from torch.utils.data import Dataset
35 from torch.utils.data import DataLoader
36 from PIL import Image
37
38 import torch.nn as nn
39 import torch.nn.functional as F
40
41 import torch.optim as optim
42
43 import torchvision.models as models
44
45 class MyDataset(Dataset):
46     def __init__(self, img_file, label_file, transform=None, idx = None):
47         self.data = pickle.load( open( img_file, 'rb' ), encoding='bytes' )
48         self.targets = np.genfromtxt(label_file, delimiter=',', skip_header=1
49 )[:,1:]
50         if idx is not None:
51             self.targets = self.targets[idx]
52             self.data = self.data[idx]
53             self.transform = transform
54             self.targets -= 5
55
56     def __len__(self):
57         return len(self.targets)
58
59     def __getitem__(self, index):
60         img, target = self.data[index], int(self.targets[index])
61         img = Image.fromarray(img.astype('uint8'), mode='L')
62
63         if self.transform is not None:

```

```

63         img = self.transform(img)
64
65     return img, target
66
67 img_transform = transforms.Compose([
68     transforms.ToTensor(),
69     transforms.Normalize((0.5,), (0.5,)))
70 ])
71
72 train_index = np.arange(50000)
73 test_index = np.arange(50000, 60000)
74 batch_size = 32 #feel free to change it
75
76 # Read image data and their label into a Dataset class
77 train_set = MyDataset('./Train.pkl', './TrainLabels.csv', transform=
    img_transform, idx=train_index)
78 train_loader = DataLoader(train_set, batch_size=batch_size, shuffle=True,
    num_workers=2)
79 test_set = MyDataset('./Train.pkl', './TrainLabels.csv', transform=
    img_transform, idx=None)
80 test_loader = DataLoader(test_set, batch_size=batch_size, shuffle=True,
    num_workers=2)
81
82 """# Notice: In case the code blocks are not in the right order, please run
    the blocks containing train() and test() functions every time after [model,
    criterion and optimizer] definition and before calling them.
83
84 # ResNet18
85 """
86
87 resnet18 = models.resnet18()
88
89 net = resnet18
90 net.conv1 = nn.Conv2d(1, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3),
    ), bias=False)
91 net.fc = nn.Linear(512, 9)
92 # if there is a available cuda device, use GPU, else, use CPU
93 device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
94 net = net.to(device)
95
96 # set criterion to cross entropy loss
97 criterion = nn.CrossEntropyLoss()
98
99 # set learning rate to 0.001
100 optimizer = optim.SGD(net.parameters(), lr=0.001)
101 # optimizer = optim.SGD(net.parameters(), lr=0.01, momentum=0.5)
102
103 epoch = 32
104 train(epoch)
105
106 test()
107
108 """# AlexNet"""
109
110 alexnet = models.alexnet()
111
112 net = alexnet
113 net.features[0] = nn.Conv2d(1, 64, kernel_size=(11, 11), stride=(4, 4),
    padding=(2, 2))
114 net.classifier[6] = nn.Linear(in_features=4096, out_features=9, bias=True)
115 # if there is a available cuda device, use GPU, else, use CPU
116 device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
117 net = net.to(device)

```

```

118
119 # set criterion to cross entropy loss
120 criterion = nn.CrossEntropyLoss()
121
122 # set learning rate to 0.001
123 # optimizer = optim.SGD(net.parameters(), lr=0.001)
124 optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.8)
125
126 epoch = 32
127 train(epoch)
128
129 test()
130
131 """# VGG16"""
132
133 vgg16 = models.vgg16()
134
135 net = vgg16
136 net.features[0] = nn.Conv2d(1, 64, kernel_size=(3, 3), stride=(1, 1), padding
=(1, 1))
137 net.classifier[6] = nn.Linear(in_features=4096, out_features=9, bias=True)
138 # if there is a available cuda device, use GPU, else, use CPU
139 device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
140 net = net.to(device)
141
142 # set criterion to cross entropy loss
143 criterion = nn.CrossEntropyLoss()
144
145 # set learning rate to 0.001
146 optimizer = optim.SGD(net.parameters(), lr=0.001)
147 # optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.8)
148
149 epoch = 32
150 train(epoch)
151
152 test()
153
154 """# VGG19"""
155
156 vgg19 = models.vgg19()
157
158 net = vgg19
159 net.features[0] = nn.Conv2d(1, 64, kernel_size=(3, 3), stride=(1, 1), padding
=(1, 1))
160 net.classifier[6] = nn.Linear(in_features=4096, out_features=9, bias=True)
161 # if there is a available cuda device, use GPU, else, use CPU
162 device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
163 net = net.to(device)
164
165 # set criterion to cross entropy loss
166 criterion = nn.CrossEntropyLoss()
167
168 # set learning rate to 0.001
169 # optimizer = optim.SGD(net.parameters(), lr=0.001)
170 optimizer = optim.SGD(net.parameters(), lr=0.005, momentum=0.5)
171
172 # with: optimizer = optim.SGD(net.parameters(), lr=0.001)
173 epoch = 32
174 train(epoch)
175
176 test()
177
178 # with: optimizer = optim.SGD(net.parameters(), lr=0.005, momentum=0.5)

```

```

179 epoch = 32
180 train(epoch)
181
182 test()
183
184 """# VGG19-bn"""
185
186 vgg19bn = models.vgg19_bn()
187
188 net = vgg19bn
189 net.features[0] = nn.Conv2d(1, 64, kernel_size=(3, 3), stride=(1, 1), padding
= (1, 1))
190 net.classifier[6] = nn.Linear(in_features=4096, out_features=9, bias=True)
191 # if there is a available cuda device, use GPU, else, use CPU
192 device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
193 net = net.to(device)
194
195 # set criterion to cross entropy loss
196 criterion = nn.CrossEntropyLoss()
197
198 # set learning rate to 0.001
199 # optimizer = optim.SGD(net.parameters(), lr=0.001)
200 optimizer = optim.SGD(net.parameters(), lr=0.005, momentum=0.5)
201
202 def train(num_epochs=2): # Feel free to change it
203     net.train()
204
205     running_loss = 0.0
206
207     # Here is a piece of code that reads data in batch.
208     # In each epoch all samples are read in batches using dataloader
209     for epoch in range(num_epochs):
210         for i, data in enumerate(train_loader):
211             img, label = data
212
213             img = img.to(device)
214             label = label.to(device)
215
216             # zero the parameter gradients
217             optimizer.zero_grad()
218
219             # forward + backward + optimize
220             outputs = net(img)
221
222             loss = criterion(outputs, label)
223             # loss = F.nll_loss(outputs, label)
224             loss.backward()
225             optimizer.step()
226
227             running_loss += loss.item()
228             if i % 320 == 319: # print every 320 mini-batches
229                 print('[%d, %5d] loss: %.3f' %
230                     (epoch + 1, i + 1, running_loss / 320))
231                 running_loss = 0.0
232
233             torch.save(net.state_dict(), '/model.pth')
234             torch.save(optimizer.state_dict(), '/optimizer.pth')
235
236     print('Finished Training')
237
238 def test():
239     net.eval()
240

```

```
241     correct = 0
242     total = 0
243
244     # calculate accuracy
245     with torch.no_grad():
246         for data in test_loader:
247             images, labels = data
248
249             images = images.to(device)
250             labels = labels.to(device)
251
252             outputs = net(images)
253             # get the index of the max output
254             _, predicted = torch.max(outputs.data, 1)
255             total += labels.size(0)
256             correct += (predicted == labels).sum().item()
257
258     print('Accuracy of the network: %d %%' % (
259         100 * correct / total))
260
261 # with: optimizer = optim.SGD(net.parameters(), lr=0.001)
262 epoch = 32
263 train(epoch)
264
265 test()
266
267 # with: optimizer = optim.SGD(net.parameters(), lr=0.005, momentum=0.5)
268 epoch = 32
269 train(epoch)
270
271 test()
```

```

1 # -*- coding: utf-8 -*-
2 """TensorFlow_Model.ipynb
3
4 Automatically generated by Colaboratory.
5
6 Original file is located at
7     https://colab.research.google.com/drive/1io3Hj9bFn56RH23DcKvgA2545zUIGxoc
8
9 <center><h1>Mini Project 3 - Convolutional Neural Network</h1>
10 <h4>The TensorFlow File.</h4></center>
11
12 <h3>Team Members:</h3>
13 <center>
14 Yi Zhu, 260716006<br>
15 Fei Peng, 260712440<br>
16 Yukai Zhang, 260710915
17 </center>
18 """
19
20 import torch
21 import torchvision
22 import tensorflow as tf
23 from sklearn.model_selection import train_test_split
24 import time
25 import pickle
26 import numpy as np
27 import pandas as pd
28 from PIL import Image
29 import torchvision.transforms as transforms
30 import matplotlib.pyplot as plt
31
32 from torch.utils.data import Dataset
33 from torch.utils.data import DataLoader
34
35 import tensorflow.keras as keras
36 from tensorflow.keras.applications.vgg19 import VGG19
37 from tensorflow.keras.models import Model
38 from tensorflow.keras.layers import Dense, GlobalAveragePooling2D, Dropout
39 from tensorflow.keras.optimizers import SGD
40
41 # Commented out IPython magic to ensure Python compatibility.
42 from google.colab import drive
43 drive.mount("/content/drive")
44
45 # %cd '/content/drive/My Drive/ECSE_551_F_2020/Mini_Project_03/'
46
47 """# Import and Preprocess Dataset"""
48
49 class MyDataset(Dataset):
50     def __init__(self, img_file, label_file, transform=None, idx = None):
51         self.data = pickle.load( open( img_file, 'rb' ), encoding='bytes' )
52         self.targets = np.genfromtxt(label_file, delimiter=',', skip_header=1
53 )[:,1:]
54         if idx is not None:
55             self.targets = self.targets[idx]
56             self.data = self.data[idx]
57             self.transform = transform
58             self.targets -= 5
59
60     def __len__(self):
61         return len(self.targets)
62
63     def __getitem__(self, index):

```

```

63         img, target = self.data[index], int(self.targets[index])
64         img = Image.fromarray(img.astype('uint8'), mode='L')
65
66         if self.transform is not None:
67             img = self.transform(img)
68
69         return img, target
70
71 img_transform = transforms.Compose([
72     transforms.ToTensor(),
73     transforms.Normalize((0.5,), (0.5,)))
74 ])
75
76 batch_size = 32 #feel free to change it
77 # Read image data and their label into a Dataset class
78 train_set = MyDataset('./Train.pkl', './TrainLabels.csv', transform=
    img_transform, idx=None)
79
80 train_set_data = np.repeat(train_set.data[..., np.newaxis], 3, -1)
81 # train_x, x_validation, train_y, y_validation = train_test_split(
    train_set_data, train_set.targets, test_size=0.20, random_state=0)
82 train_x = train_set_data
83 train_y = train_set.targets
84 print(train_x.shape, train_y.shape)
85
86 """# Define Model"""
87
88 vgg19 = VGG19(weights=None, include_top=False)
89
90 # add a global spatial average pooling layer
91 x = vgg19.output
92 x = GlobalAveragePooling2D()(x)
93 x = Dropout(0.3)(x)
94 x = Dense(128, activation='relu')(x)
95 predictions = Dense(9, activation='softmax')(x)
96
97 # this is the model to train
98 model = Model(inputs=vgg19.input, outputs=predictions)
99
100 # make all layers trainable
101 for layer in vgg19.layers:
102     layer.trainable = True
103
104 """# Train"""
105
106 model.compile(keras.optimizers.SGD(learning_rate=0.001, momentum=0.7), loss='
    sparse_categorical_crossentropy', metrics=['accuracy'])
107
108 # model.fit(train_x, train_y, epochs=45, validation_data=(x_validation,
    y_validation), batch_size=32)
109 model.fit(train_x, train_y, epochs=65, batch_size=32)
110
111 """# Output Export"""
112
113 test_data = pickle.load(open( './Test.pkl', 'rb' ), encoding='bytes')
114 test_data = np.repeat(test_data[..., np.newaxis], 3, -1)
115
116 from google.colab import files
117 import pandas as pd
118
119 y_test = model.predict(test_data, batch_size=32)
120 X_id = np.arange(y_test.shape[0])
121

```

```
122 y_class = np.argmax(y_test, axis=1) + 5
123 print(y_class)
124 result = {'id': X_id, 'class': y_class}
125
126 df = pd.DataFrame(data=result)
127 df.to_csv('result.csv', index=False)
128 files.download('result.csv')
```